

# Rasterization & The Graphics Pipeline

Fast Approximations for Real-Time Graphics

---

Ashrafur Rahman

Adjunct Lecturer

Department of Computer Science and Engineering  
Bangladesh University of Engineering and Technology (BUET)

## Introduction

Philosophy

The Speed vs Accuracy Trade-off

Realtime Challenge

## The GPU Evolution

History

GPU Domination

Modern GPU Architecture

## The Modern Graphics Pipeline

Definition

Stages

Advanced Pipeline

Data Flow Through the Pipeline

Pipeline Stage Types

# Introduction

---

# Why Rasterization?

# Why Rasterization?



**Valorant - 120 FPS Gaming**

# Why Rasterization?



**Valorant** - 120 FPS Gaming



**Up** - 30 hours per frame

- **Real-time constraint:** Games need 60-120 FPS

# Why Rasterization?



**Valorant** - 120 FPS Gaming



**Up** - 30 hours per frame

- **Real-time constraint:** Games need 60-120 FPS
- **Interactive experience:** User input must feel responsive

# Why Rasterization?



**Valorant** - 120 FPS Gaming



**Up** - 30 hours per frame

- **Real-time constraint:** Games need 60-120 FPS
- **Interactive experience:** User input must feel responsive
- **Trade-off:** Sacrifice physical accuracy for speed



# Why Rasterization?



**Valorant** - 120 FPS Gaming

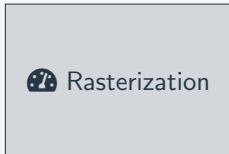


**Up** - 30 hours per frame

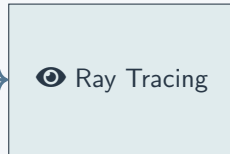
- **Real-time constraint:** Games need 60-120 FPS
- **Interactive experience:** User input must feel responsive
- **Trade-off:** Sacrifice physical accuracy for speed
- **Goal:** Images that look good enough, delivered fast enough

# Rasterization vs Ray Tracing: The Fundamental Choice

## Fast Approximations



## Physical Accuracy



Trade-off

### Rasterization:

- 60-240 FPS
- Clever approximations
- Hardware optimized
- "Good enough" quality

### Ray Tracing:

- $\approx 0$  FPS
- Physical simulation
- Computationally heavy
- Photorealistic

# The Real-Time Graphics Challenge

## Time Budget at 60 FPS

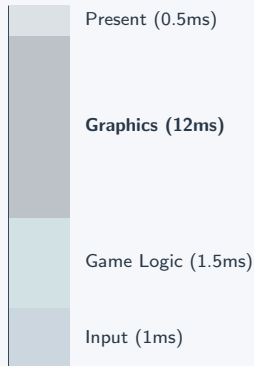
$$\frac{1}{60} = 16.67 \text{ milliseconds per frame}$$

### What needs to happen:

- Process input
- Update game logic
- Render graphics
- Present to screen

**Graphics budget:** ~10-12ms

16.67ms

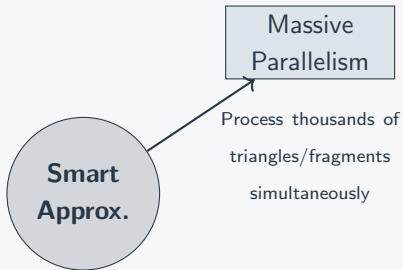


# How Rasterization Achieves Speed

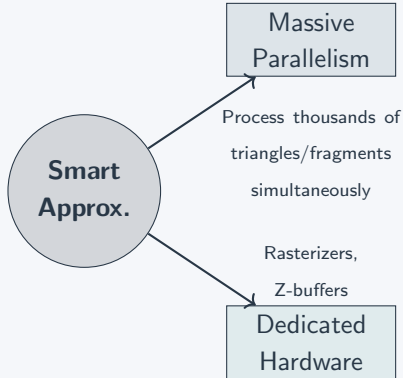


**Smart  
Approx.**

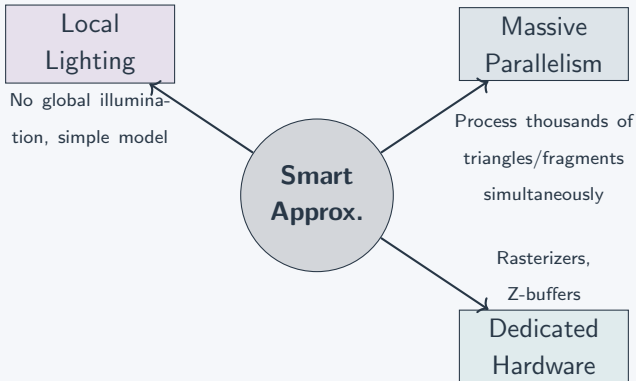
# How Rasterization Achieves Speed



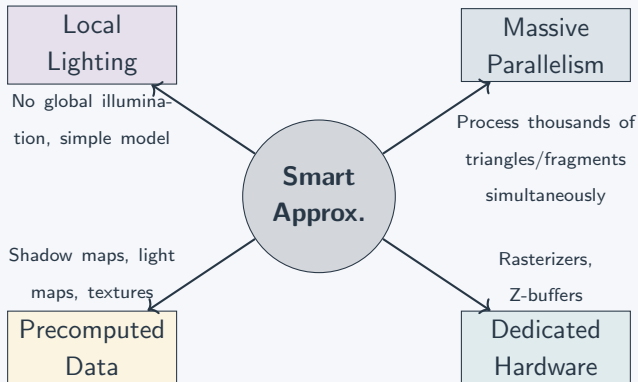
# How Rasterization Achieves Speed



# How Rasterization Achieves Speed



# How Rasterization Achieves Speed





# The Clever Approximations

## What We Skip

- **Global illumination:**  
No light bouncing
- **Perfect shadows:** Use shadow maps
- **Perfect reflections:**  
Use environment maps
- **Complex materials:**  
Simplified BRDFs

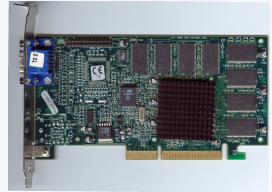
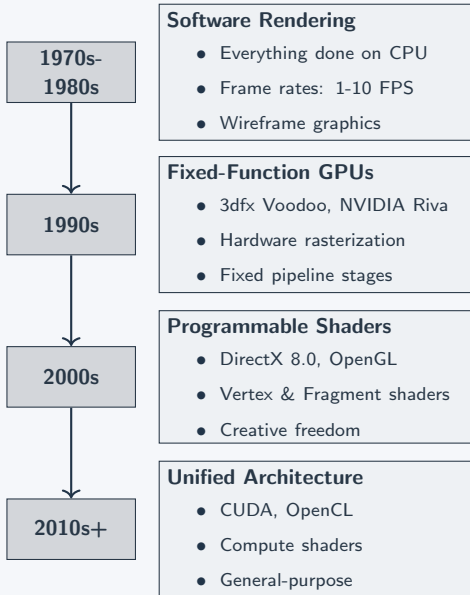
## What We Gain

- **Predictable performance:** Linear with triangle count
- **Hardware optimization:**  
Purpose-built silicon
- **Real-time interaction:**  
Immediate feedback
- **Scalable quality:**  
Adjust for performance

# The GPU Evolution

---

# A Brief History



3dfx Voodoo 3 - 1999



NVIDIA GeForce 5090 - 2025

# Why GPUs Dominate Graphics

## CPU

4-16 complex cores

Large caches

Branch prediction

Out-of-order execution

# Why GPUs Dominate Graphics

## CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

## GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

# Why GPUs Dominate Graphics

## CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

## GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

## Serial Tasks

- Complex logic
- Branching
- Low latency

# Why GPUs Dominate Graphics

## CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

## Serial Tasks

- Complex logic
- Branching
- Low latency

## GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

## Parallel Tasks

- Simple operations
- Same instruction
- High throughput

# Why GPUs Dominate Graphics

## CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

## Serial Tasks

- Complex logic
- Branching
- Low latency

## GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

## Parallel Tasks

- Simple operations
- Same instruction
- High throughput

## Perfect Match: Graphics + GPU

**Graphics pipeline stages** process thousands of vertices/fragments *independently*

⇒ Ideal for massively parallel GPU architecture



# Modern GPU: The Graphics Powerhouse

## Hardware Implementation


### GPU handles entire pipeline:

- **Vertex processing:** Shader cores
- **Rasterization:** Fixed-function units
- **Fragment processing:** Shader cores
- **Memory operations:** ROPs

### GPU Driver handles:

- Command submission
- State management
- Resource allocation

## GPU Chip



# The Modern Graphics Pipeline

---

# Why a Graphics Pipeline?

## The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)

# Why a Graphics Pipeline?

## The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**

# Why a Graphics Pipeline?

## The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**

# Why a Graphics Pipeline?

## The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**
- Each stage can be **efficiently** implemented in hardware

# Why a Graphics Pipeline?

## The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**
- Each stage can be **efficiently** implemented in hardware
- Stages can be **parallelized** for high throughput

# Why a Graphics Pipeline?

## The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**
- Each stage can be **efficiently** implemented in hardware
- Stages can be **parallelized** for high throughput

## The Graphics Pipeline

The graphics pipeline is a sequence of stages that process vertices and fragments in parallel, transforming 3D models into 2D images.




# Pipeline Stages at a Glance

Input Assembly

**Input Assembly:** Pull vertex data (positions, normals, UVs) into the pipeline.

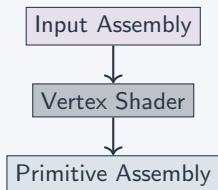
# Pipeline Stages at a Glance



**Input Assembly:** Pull vertex data (positions, normals, UVs) into the pipeline.

**Vertex Shader:** Programmable stage — transform each vertex from model to clip space.

# Pipeline Stages at a Glance

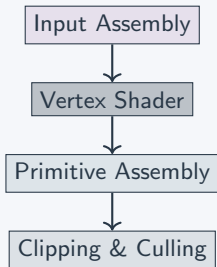


**Input Assembly:** Pull vertex data (positions, normals, UVs) into the pipeline.

**Vertex Shader:** Programmable stage — transform each vertex from model to clip space.

**Primitive Assembly:** Group vertices into triangles (or other primitives).

# Pipeline Stages at a Glance



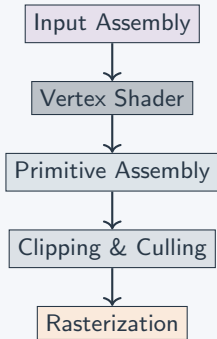
**Input Assembly:** Pull vertex data (positions, normals, UVs) into the pipeline.

**Vertex Shader:** Programmable stage — transform each vertex from model to clip space.

**Primitive Assembly:** Group vertices into triangles (or other primitives).

**Clipping & Culling:** Discard or trim primitives outside the view frustum.

# Pipeline Stages at a Glance



**Input Assembly:** Pull vertex data (positions, normals, UVs) into the pipeline.


**Vertex Shader:** Programmable stage — transform each vertex from model to clip space.

**Primitive Assembly:** Group vertices into triangles (or other primitives).

**Clipping & Culling:** Discard or trim primitives outside the view frustum.

**Rasterization:** Convert triangles into a grid of fragments (potential pixels).

# Pipeline Stages at a Glance



**Input Assembly:** Pull vertex data (positions, normals, UVs) into the pipeline.

**Vertex Shader:** Programmable stage — transform each vertex from model to clip space.

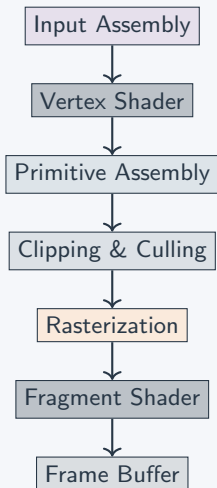
**Primitive Assembly:** Group vertices into triangles (or other primitives).

**Clipping & Culling:** Discard or trim primitives outside the view frustum.

**Rasterization:** Convert triangles into a grid of fragments (potential pixels).

**Fragment Shader:** Programmable stage — compute final color of each fragment.

# Pipeline Stages at a Glance



**Input Assembly:** Pull vertex data (positions, normals, UVs) into the pipeline.

**Vertex Shader:** Programmable stage — transform each vertex from model to clip space.

**Primitive Assembly:** Group vertices into triangles (or other primitives).

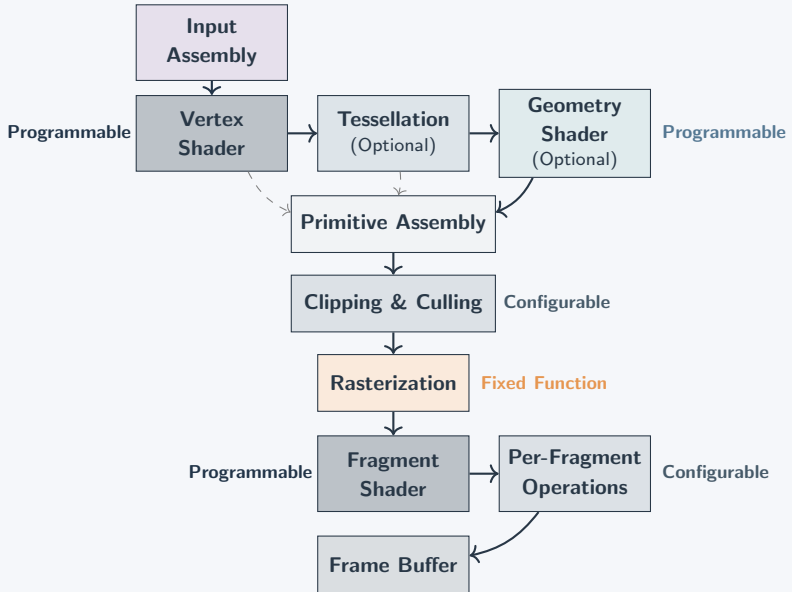
**Clipping & Culling:** Discard or trim primitives outside the view frustum.

**Rasterization:** Convert triangles into a grid of fragments (potential pixels).

**Fragment Shader:** Programmable stage — compute final color of each fragment.

**Frame Buffer:** Blend, depth-test, and write pixels to the screen.

# Modern Advanced Pipeline





# Data Flow Through the Pipeline

**Input Data**



3D Vertices

# Data Flow Through the Pipeline

**Input Data**



3D Vertices




**Vertex Stage**




Screen Space  
(Unclipped)


# Data Flow Through the Pipeline



# Data Flow Through the Pipeline



# Data Flow Through the Pipeline



# Programmable vs Fixed Function Stages

## Programmable Stages

**You write the code:**

- **Vertex Shader:**  
Transform positions,  
compute lighting
- **Tessellation:** Subdivide  
surfaces adaptively
- **Geometry Shader:**  
Generate/modify  
primitives
- **Fragment Shader:**  
Compute final pixel colors

**Maximum flexibility**

# Programmable vs Fixed Function Stages

## Programmable Stages

**You write the code:**

- **Vertex Shader:**  
Transform positions, compute lighting
- **Tessellation:** Subdivide surfaces adaptively
- **Geometry Shader:**  
Generate/modify primitives
- **Fragment Shader:**  
Compute final pixel colors

**Maximum flexibility**


## Fixed Function Stages

**Hardware handles it:**

- **Primitive Assembly:**  
Group vertices into triangles
- **Clipping:** Remove off-screen geometry
- **Rasterization:** Convert triangles to pixels
- **Depth Testing:** Z-buffer comparisons






**Maximum performance**

# Questions?





## References & Further Reading

-  Peter Shirley and Steve Marschner et al. *Fundamentals of Computer Graphics (4th Edition)*. CRC Press, 2016.  
Available as PDF
-  Matt Pharr, Wenzel Jakob, and Greg Humphreys. *Physically Based Rendering: From Theory to Implementation (4th Edition)*. Morgan Kaufmann, 2023.  
Available online
-  Peter Shirley. *Ray Tracing in One Weekend*. Self-published, 2016–2020.  
Project Website
-  MIT OpenCourseWare: 6.837 Computer Graphics.  
[ocw.mit.edu/6-837](https://ocw.mit.edu/6-837)
-  Scratchapixel: Learn Computer Graphics Programming.  
[scratchapixel.com](https://scratchapixel.com)