

Rasterization & The Graphics Pipeline

Fast Approximations for Real-Time Graphics

Ashrafur Rahman

Adjunct Lecturer

Department of Computer Science and Engineering
Bangladesh University of Engineering and Technology (BUET)

Index

Introduction

Philosophy

The Speed vs Accuracy Trade-off

Realtime Challenge

The GPU Evolution

History

GPU Domination

Modern GPU Architecture

The Modern Graphics Pipeline

Definition

Stages

Advanced Pipeline

Data Flow Through the Pipeline

Pipeline Stage Types

APIs & Shading Languages

Graphics APIs

Shading Languages

GLSL Overview

Input Assembly

Overview

Introduction

Why Rasterization?

Why Rasterization?



Valorant - 120 FPS Gaming

Why Rasterization?



Valorant - 120 FPS Gaming



Up - 30 hours per frame

- **Real-time constraint:** Games need 60-120 FPS

Why Rasterization?



Valorant - 120 FPS Gaming



Up - 30 hours per frame

- **Real-time constraint:** Games need 60-120 FPS
- **Interactive experience:** User input must feel responsive

Why Rasterization?



Valorant - 120 FPS Gaming



Up - 30 hours per frame

- **Real-time constraint:** Games need 60-120 FPS
- **Interactive experience:** User input must feel responsive
- **Trade-off:** Sacrifice physical accuracy for speed

Why Rasterization?



Valorant - 120 FPS Gaming

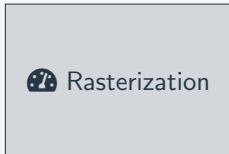


Up - 30 hours per frame

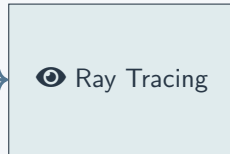
- **Real-time constraint:** Games need 60-120 FPS
- **Interactive experience:** User input must feel responsive
- **Trade-off:** Sacrifice physical accuracy for speed
- **Goal:** Images that look good enough, delivered fast enough

Rasterization vs Ray Tracing: The Fundamental Choice

Fast Approximations



Physical Accuracy



Trade-off

Rasterization:

- 60-240 FPS
- Clever approximations
- Hardware optimized
- "Good enough" quality

Ray Tracing:

- ≈ 0 FPS
- Physical simulation
- Computationally heavy
- Photorealistic

The Real-Time Graphics Challenge

Time Budget at 60 FPS

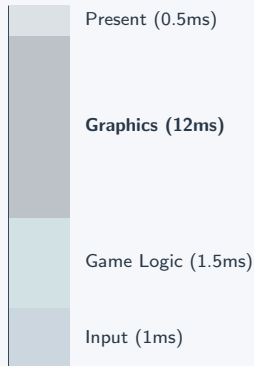
$$\frac{1}{60} = 16.67 \text{ milliseconds per frame}$$

What needs to happen:

- Process input
- Update game logic
- Render graphics
- Present to screen

Graphics budget: ~10-12ms

16.67ms

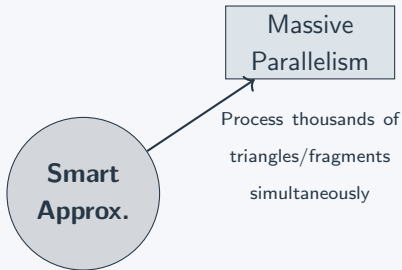


How Rasterization Achieves Speed

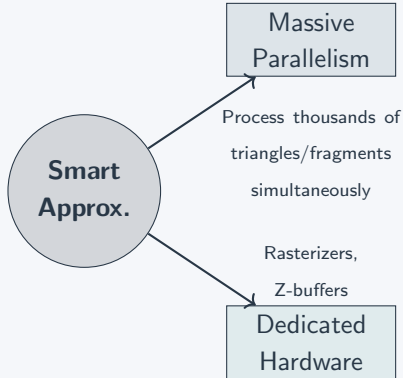


**Smart
Approx.**


How Rasterization Achieves Speed



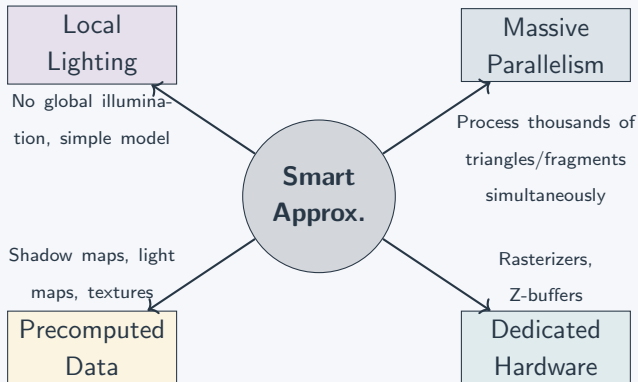
How Rasterization Achieves Speed



How Rasterization Achieves Speed



How Rasterization Achieves Speed



The Clever Approximations

What We Skip

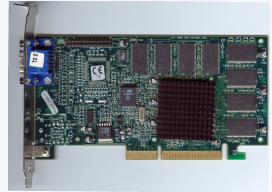
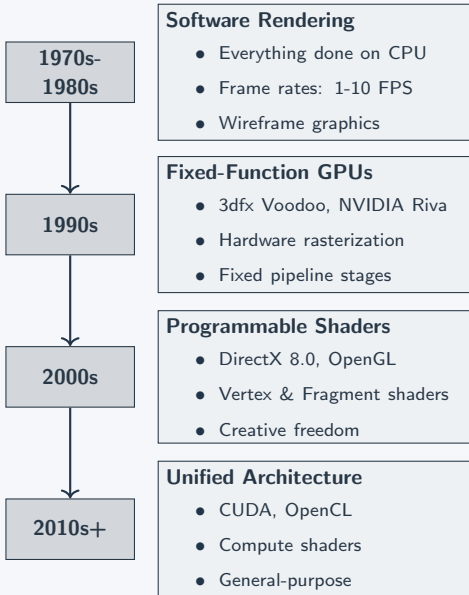
- **Global illumination:**
No light bouncing
- **Perfect shadows:** Use shadow maps
- **Perfect reflections:**
Use environment maps
- **Complex materials:**
Simplified BRDFs

What We Gain

- **Predictable performance:** Linear with triangle count
- **Hardware optimization:**
Purpose-built silicon
- **Real-time interaction:**
Immediate feedback
- **Scalable quality:**
Adjust for performance

The GPU Evolution

A Brief History



3dfx Voodoo 3 - 1999



NVIDIA GeForce 5090 - 2025

Why GPUs Dominate Graphics

CPU

4-16 complex cores

Large caches

Branch prediction

Out-of-order execution

Why GPUs Dominate Graphics

CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

Why GPUs Dominate Graphics

CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

Serial Tasks

- Complex logic
- Branching
- Low latency

Why GPUs Dominate Graphics

CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

Serial Tasks

- Complex logic
- Branching
- Low latency

GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

Parallel Tasks

- Simple operations
- Same instruction
- High throughput

Why GPUs Dominate Graphics

CPU

- 4-16 complex cores
- Large caches
- Branch prediction
- Out-of-order execution

Serial Tasks

- Complex logic
- Branching
- Low latency

GPU

- 1000s of simple cores
- Small caches
- SIMD execution
- Throughput optimized

Parallel Tasks

- Simple operations
- Same instruction
- High throughput

Perfect Match: Graphics + GPU

Graphics pipeline stages process thousands of vertices/fragments *independently*

⇒ Ideal for massively parallel GPU architecture

Modern GPU: The Graphics Powerhouse

Hardware Implementation


GPU handles entire pipeline:

- **Vertex processing:** Shader cores
- **Rasterization:** Fixed-function units
- **Fragment processing:** Shader cores
- **Memory operations:** ROPs

GPU Driver handles:

- Command submission
- State management
- Resource allocation

GPU Chip



The Modern Graphics Pipeline

Why a Graphics Pipeline?

The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)

Why a Graphics Pipeline?

The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**

Why a Graphics Pipeline?

The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**

Why a Graphics Pipeline?

The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**
- Each stage can be **efficiently** implemented in hardware

Why a Graphics Pipeline?

The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**
- Each stage can be **efficiently** implemented in hardware
- Stages can be **parallelized** for high throughput

Why a Graphics Pipeline?

The Rasterization Process

- GPU process vast numbers of **vertices** and **pixels** every frame (millions per second)
- Each undergoes a **series of identical operations**
- The operations can be divided into **stages**
- Each stage can be **efficiently** implemented in hardware
- Stages can be **parallelized** for high throughput

The Graphics Pipeline


The graphics pipeline is a sequence of stages that process vertices and fragments in parallel, transforming 3D models into 2D images.

Pipeline Stages at a Glance

Input Assembly

Input Assembly: Pull vertex data (positions, normals, UVs) into the pipeline.

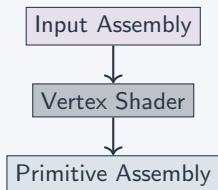
Pipeline Stages at a Glance



Input Assembly: Pull vertex data (positions, normals, UVs) into the pipeline.

Vertex Shader: Programmable stage — transform each vertex from model to clip space.

Pipeline Stages at a Glance

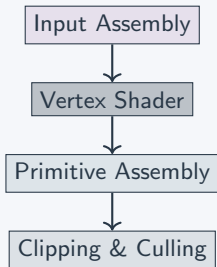


Input Assembly: Pull vertex data (positions, normals, UVs) into the pipeline.

Vertex Shader: Programmable stage — transform each vertex from model to clip space.

Primitive Assembly: Group vertices into triangles (or other primitives).

Pipeline Stages at a Glance



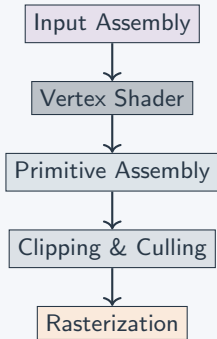
Input Assembly: Pull vertex data (positions, normals, UVs) into the pipeline.

Vertex Shader: Programmable stage — transform each vertex from model to clip space.

Primitive Assembly: Group vertices into triangles (or other primitives).

Clipping & Culling: Discard or trim primitives outside the view frustum.

Pipeline Stages at a Glance



Input Assembly: Pull vertex data (positions, normals, UVs) into the pipeline.

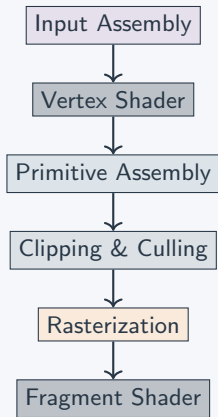
Vertex Shader: Programmable stage — transform each vertex from model to clip space.

Primitive Assembly: Group vertices into triangles (or other primitives).

Clipping & Culling: Discard or trim primitives outside the view frustum.

Rasterization: Convert triangles into a grid of fragments (potential pixels).

Pipeline Stages at a Glance



Input Assembly: Pull vertex data (positions, normals, UVs) into the pipeline.

Vertex Shader: Programmable stage — transform each vertex from model to clip space.

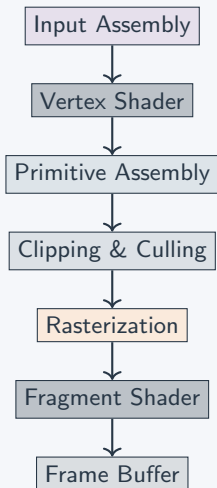
Primitive Assembly: Group vertices into triangles (or other primitives).

Clipping & Culling: Discard or trim primitives outside the view frustum.

Rasterization: Convert triangles into a grid of fragments (potential pixels).

Fragment Shader: Programmable stage — compute final color of each fragment.

Pipeline Stages at a Glance



Input Assembly: Pull vertex data (positions, normals, UVs) into the pipeline.

Vertex Shader: Programmable stage — transform each vertex from model to clip space.

Primitive Assembly: Group vertices into triangles (or other primitives).

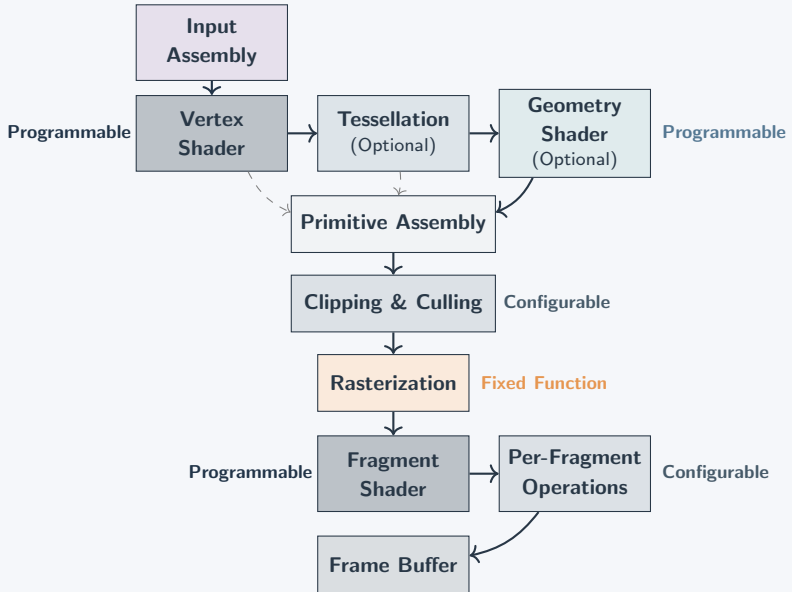
Clipping & Culling: Discard or trim primitives outside the view frustum.

Rasterization: Convert triangles into a grid of fragments (potential pixels).

Fragment Shader: Programmable stage — compute final color of each fragment.

Frame Buffer: Blend, depth-test, and write pixels to the screen.

Modern Advanced Pipeline



Data Flow Through the Pipeline

Input Data



3D Vertices

Data Flow Through the Pipeline

Input Data



3D Vertices




Vertex Stage

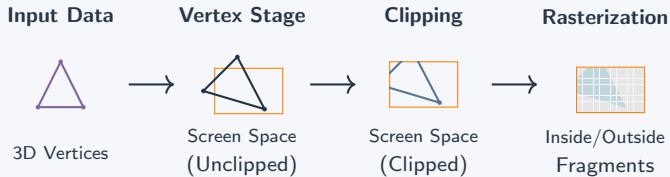


Screen Space
(Unclipped)

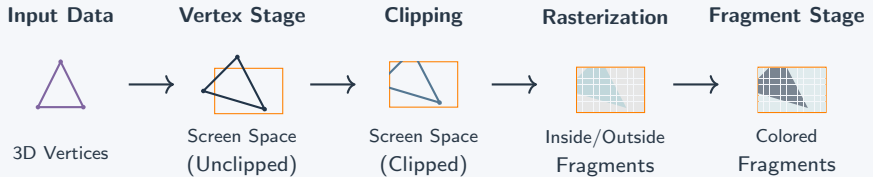
Data Flow Through the Pipeline



Data Flow Through the Pipeline



Data Flow Through the Pipeline



Programmable vs Fixed Function Stages

Programmable Stages

You write the code:

- **Vertex Shader:**
Transform positions,
compute lighting
- **Tessellation:** Subdivide
surfaces adaptively
- **Geometry Shader:**
Generate/modify
primitives
- **Fragment Shader:**
Compute final pixel colors

Maximum flexibility

Programmable vs Fixed Function Stages

Programmable Stages

You write the code:

- **Vertex Shader:**
Transform positions,
compute lighting
- **Tessellation:** Subdivide
surfaces adaptively
- **Geometry Shader:**
Generate/modify
primitives
- **Fragment Shader:**
Compute final pixel colors

Maximum flexibility

Fixed Function Stages

Hardware handles it:

- **Primitive Assembly:**
Group vertices into
triangles
- **Clipping:** Remove
off-screen geometry
- **Rasterization:** Convert
triangles to pixels
- **Depth Testing:** Z-buffer
comparisons

Maximum performance

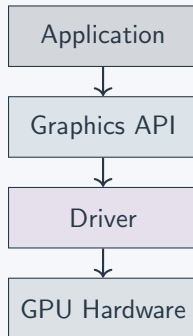
APIs & Shading Languages

Graphics APIs Overview

What is a Graphics API?

An **Application Programming Interface** that provides:




- Commands to control the GPU
- Abstraction over hardware differences
- Standard interface for graphics operations



OpenGL

Open Graphics Library

Perhaps the most widely used and most beginner-friendly graphics API. OpenGL was designed to be a cross-platform standard for rendering 2D and 3D graphics.





- Stable APIs
- High-level abstraction
- , ,  Support
- Two modes -
 - **Immediate Mode** - Deprecated, fixed function (Used by iGraphics)
 - **Retained Mode** - Modern OpenGL, uses shaders



Vulkan

Low-overhead, Cross-platform Graphics API

Developed by the Khronos Group as a modern successor to OpenGL. Designed for high-performance, multi-threaded rendering.

- Low-level control over GPU
- Better CPU-GPU parallelism
- Explicit memory and resource management
- , ,  Support
-  Support via MoltenVK





Major Graphics APIs

DirectX (Direct3D)

Microsoft's Graphics API for Windows and Xbox

A powerful API suite used primarily for game development on Windows platforms.

- Direct3D for 3D rendering
- Deep integration with Windows OS and drivers
- High performance with hardware vendor optimizations
-   Support only



Metal

Apple's Low-level Graphics API

Designed to maximize performance on Apple devices, replacing OpenGL on Apple platforms.

- Low-overhead, low-level access
- Unified graphics and compute
- Tight integration with Apple hardware
- 🍏 Support only (macOS, iOS, iPadOS)



Purpose

Shading languages allow programmers to write code that runs on the GPU for:

- **Vertex processing**
(transformations)
- **Fragment processing**
(lighting, texturing, effects)
- **Compute operations**
(general-purpose GPU computing)

Major Shading Languages:

- **GLSL** (OpenGL)
- **HLSL** (DirectX)
- **MSL** (Metal)
- **SPIR-V** (Vulkan)
Can be compiled from
GLSL or HLSL

GLSL: OpenGL Shading Language

GLSL Characteristics

- C-like syntax
- Built-in vector/matrix types
- Version-specific features

Data Types

- float, int, bool
- vec2, vec3, vec4
- mat2, mat3, mat4
- sampler2D, samplerCube

```
#version 330 core
// Vertex shader inputs
layout (location = 0) in vec3 aPos;
layout (location = 1) in vec3 aNormal;
layout (location = 2) in vec2 aTexCoord;
// Outputs to fragment shader
out vec3 FragPos;
out vec3 Normal;
out vec2 TexCoord;
// Uniform variables
uniform mat4 model;
uniform mat4 view;
uniform mat4 projection;

void main() {
    FragPos = vec3(model * vec4(aPos, 1.0));
    Normal = mat3(transpose(inverse(model)))
               * aNormal;
    TexCoord = aTexCoord;

    gl_Position = projection
                   * view
                   * vec4(FragPos, 1.0);
}
```

Storage Qualifiers

- `in` - Input from previous stage
- `out` - Output to next stage
- `uniform` - Buffer from CPU

Layout Qualifiers

Used to explicitly assign indices or binding points to resources.

- `location` - Attribute/output index
- `binding` - Texture/uniform buffer slot

```
#version 330 core
// Fragment shader

in vec3 FragPos;
in vec3 Normal;
in vec2 TexCoord;
out vec4 FragColor;

uniform vec3 lightPos;
uniform vec3 viewPos;
uniform sampler2D texture_diffuse1;

void main() {
    vec3 color = texture(texture_diffuse1,
                          TexCoord).rgb;

    // Ambient
    vec3 ambient = 0.1 * color;

    // Diffuse
    vec3 norm = normalize(Normal);
    vec3 lightDir = normalize(
        lightPos - FragPos);
    float diff = max(dot(norm, lightDir),
                     0.0);
    vec3 diffuse = diff * color;

    // Result
    vec3 result = ambient + diffuse;
    FragColor = vec4(result, 1.0);
}
```


Input Assembly

Input Assembly Stage

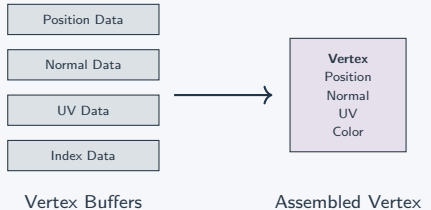
Input Assembly

Input: Vertex data from application

Output: Organized vertex streams for vertex shader

Purpose:

- Pull vertex data from memory
- Organize data into vertex attributes
- Handle indexed vs non-indexed drawing
- Set up primitive topology



Vertex Attributes

Common Vertex Attributes

- **Position:** 3D coordinates (vec3)
- **Normal:** Surface normal vector (vec3)
- **Texture Coordinates:** UV mapping (vec2)
- **Color:** Vertex color (vec3/vec4)

Vertex Layout Example

$$\text{Vertex} = \begin{cases} \text{Position:} & (x, y, z) \\ \text{Normal:} & (n_x, n_y, n_z) \\ \text{UV:} & (u, v) \\ \text{Color:} & (r, g, b, a) \end{cases} \quad (1)$$

Total size: $3 + 3 + 2 + 4 = 12$ floats = 48 bytes

Position

Normal

UV

Color

12 bytes

12 bytes

8 bytes

16 bytes

Primitive Topology

Primitive Types

Points: Individual vertices

- Used for particle systems
- Point sprites

Lines: Connected line segments

- Wireframe rendering
- Debug visualization

Triangles: Most common primitive

- Standard for 3D surfaces
- Hardware optimized

Points



Lines



Triangles



Indexed vs Non-Indexed Drawing

Non-Indexed Drawing

Direct vertex specification

Each vertex is specified multiple times for shared vertices.

Problem:

- Vertex duplication
- Increased memory usage
- Inefficient for complex meshes

Indexed Drawing

Vertices referenced by indices

Each vertex is stored once, and indices specify how to connect them.

Benefits:

- No vertex duplication
- Lower memory usage
- Vertex cache friendly

Non-Indexed

Vertices: A, B, C, B, C, D

33% less data



Indexed

Vertices: A, B, C, D

Indices: 0, 1, 2, 1, 2, 3

Vertex Shader

Vertex Shader Stage

Vertex Shader

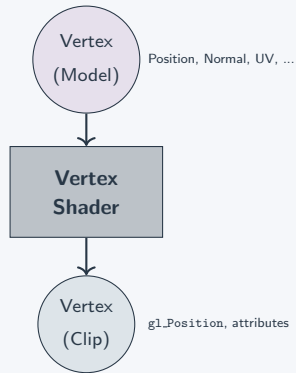
Input: Individual vertices with attributes

Output: Transformed vertices in clip space

Purpose:

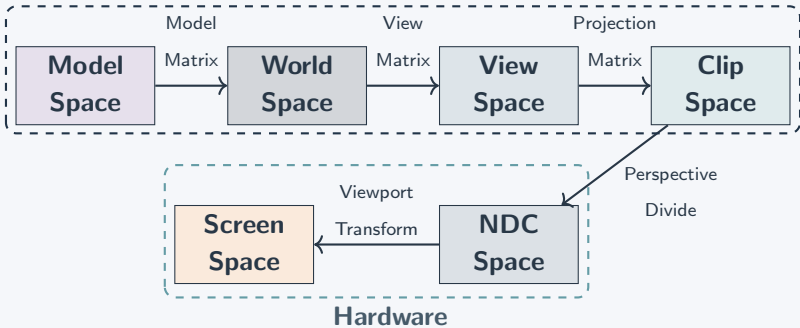
- Transform vertex positions through coordinate spaces
- Pass attributes to next stage
- Optional - Apply animations and deformations
- Optional - Calculate per-vertex lighting (Gouraud shading)

Programmable stage - you write the code!



The Transformation Pipeline

Vertex Shader



Transformation Matrices

Model Matrix

Purpose: Object-to-world transformation

$$\mathbf{M} = \mathbf{T} \cdot \mathbf{R} \cdot \mathbf{S}$$

- **T**: Translation
- **R**: Rotation
- **S**: Scale

Transforms from model's local coordinates to world coordinates.

View Matrix

Purpose: World-to-camera transformation

$$\mathbf{V} = \text{lookAt}(\text{eye}, \text{target}, \text{up})$$

Transforms from world coordinates to camera/eye coordinates.

Camera is at origin, looking down -Z axis.

Transformation Matrices

Projection Matrix

Purpose: Camera-to-clip space transformation

Perspective:

$$\mathbf{P}_{\text{persp}} = \begin{pmatrix} \frac{1}{\tan(\text{fov}/2) \cdot \text{aspect}} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(\text{fov}/2)} & 0 & 0 \\ 0 & 0 & \frac{f+n}{n-f} & \frac{2fn}{n-f} \\ 0 & 0 & -1 & 0 \end{pmatrix}$$

Orthographic:

$$\mathbf{P}_{\text{ortho}} = \begin{pmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{-2}{f-n} & -\frac{f+n}{f-n} \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

Clipping and Perspective Division

Clipping Phase

Purpose: Remove primitives outside the viewing frustum

- Operates in **clip space** (homogeneous coordinates)
- Keeps only geometry inside the cube: $-w \leq x, y, z \leq w$
- Output goes through **perspective division**:

$$\text{NDC} = \left(\frac{x}{w}, \frac{y}{w}, \frac{z}{w} \right)$$

Normalized Device Coordinates (NDC)

Canonical cube from $(-1, -1, -1)$ to $(1, 1, 1)$

- Independent of screen resolution
- Defines final visible region before pixel conversion
- x : left/right, y : bottom/top, z : near/far

Viewport Transformation and Screen Space

Viewport Transformation

Maps NDC coordinates to screen pixels:

$$\text{screen}_x = \left(\frac{\text{ndc}_x + 1}{2} \right) \cdot \text{width} + x_{\text{offset}}$$

$$\text{screen}_y = \left(\frac{\text{ndc}_y + 1}{2} \right) \cdot \text{height} + y_{\text{offset}}$$


- Converts from $[-1, 1]$ range to pixel coordinates
- Accounts for viewport size and position

Screen Space

Final 2D position used for rasterization

- Measured in pixels
- Used by rasterizer to generate fragments

Frustum → NDC → Screen Space



Vertex Shader Example

```
#version 330 core
// Vertex shader inputs
layout (location = 0) in vec3 aPos;
layout (location = 1) in vec3 aNormal;
layout (location = 2) in vec2 aTexCoord;
// Outputs to fragment shader
out vec3 FragPos;
out vec3 Normal;
out vec2 TexCoord;
// Uniform variables
uniform mat4 model;
uniform mat4 view;
uniform mat4 projection;

void main() {
    FragPos = vec3(model * vec4(aPos, 1.0));
    Normal = mat3(transpose(inverse(model)))
              * aNormal;
    TexCoord = aTexCoord;

    gl_Position = projection
                  * view
                  * vec4(FragPos, 1.0);
}
```

Output Variables

gl_Position -
Clip space position

Gouraud Shading

Gouraud (Per-Vertex) Shading


Compute lighting at vertices, interpolate across triangles

Process:

1. Calculate lighting at each vertex
2. Output vertex color
3. Hardware interpolates colors across triangle

Pros: Fast, good for distant objects

Cons: Poor specular highlights, faceted appearance



Flat




Gouraud








Phong

Questions?



References & Further Reading

-  Peter Shirley and Steve Marschner et al. *Fundamentals of Computer Graphics (4th Edition)*. CRC Press, 2016.
Available as PDF
-  Matt Pharr, Wenzel Jakob, and Greg Humphreys. *Physically Based Rendering: From Theory to Implementation (4th Edition)*. Morgan Kaufmann, 2023.
Available online
-  Peter Shirley. *Ray Tracing in One Weekend*. Self-published, 2016–2020.
Project Website
-  MIT OpenCourseWare: 6.837 Computer Graphics.
ocw.mit.edu/6-837
-  Scratchapixel: Learn Computer Graphics Programming.
scratchapixel.com